


# High Performance Computing


Term 4 2018/2019

Lecture 9

# Step complexity vs work complexity




# Step complexity vs work complexity



Step complexity = 3  
Work complexity = 7

# Serial reduce

$$\boxed{1} + \boxed{2} + \boxed{3} + \boxed{4} + \boxed{5}$$




Reduce:

- 1) set of elements
- 2) reduction operator (binary and associative)

step complexity: 4  
work complexity: 4

# Parallel reduce

$$\boxed{1} + \boxed{2} + \boxed{3} + \boxed{4} + \boxed{5}$$




step complexity: 2

work complexity: 3

number of elements: 4

# Parallel reduce

$$\boxed{1} + \boxed{2} + \boxed{3} + \boxed{4} + \boxed{5}$$



step complexity: 3  
work complexity: 7  
number of elements: 8

$\log_2 n$  steps

# CUDA parallel reduce (naive)

```
__global__ void reduce(float * d_out, float * d_in)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int tid = threadIdx.x;

    unsigned int s = blockDim.x / 2;

    while(s>0)
    {
        if (tid<s)
        {
            d_in[myId]+=d_in[myId+s];
        }
        __syncthreads();
        s=(unsigned int)s/2;
    }

    if (tid == 0)
    {
        d_out[blockIdx.x] = d_in[myId];
        printf("%d\t%f\n", blockIdx.x, d_in[myId]);    // try without printf to see the speedup
    }
}
```

# CUDA parallel reduce (shared mem)

```
__global__ void shmem_reduce(float * d_out, float * d_in)
{
    extern float * sdata[];

    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int tid = threadIdx.x;

    sdata[tid]=d_in[myId];

    unsigned int s = blockDim.x / 2;

    while(s>0)
    {
        if (tid<s)
        {
            sdata[tid]+=sdata[tid+s];
        }
        __syncthreads();
        s=(unsigned int)s/2;
    }

    if (tid == 0)
    {
        d_out[blockIdx.x] =sdata[0];
    }
}
```

kernel call:

```
reduce<<<1024,1024,1024*sizeof(float)>>>(d_intermediate, d_in);
```



# Scan



cumulative sum (inclusive scan):

1, 3, 6, 10, 15,...


exclusive scan:

0, 1, 3, 6, 10,...

you need a binary operator and identity element (I)

# Parallel Scan using reduce

In: [1, 2, 3, 4, 5, 6, 7, 8]




Out: [1, 3, 6, 10, 15, 21, 28, 36]

step complexity:  $\log(n)$

work complexity:  $O(n^2)$

# Parallel Scan Hillis-Steele

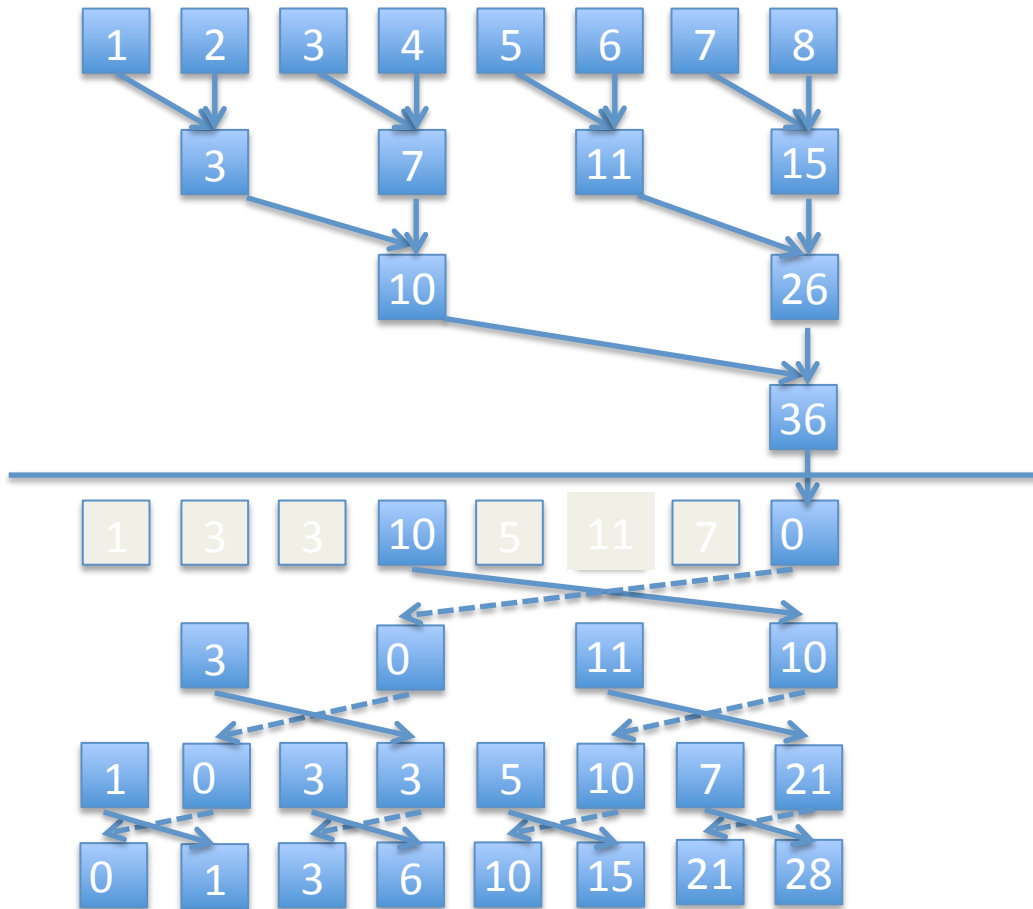
- Note that a implementation of the algorithm shown in picture requires two buffers of length  $n$  (shown is the case  $n=8=2^3$ )
- Assumption: the number  $n$  of elements is a power of 2:  $n=2^M$



steps:  $\log(n)$

work:  $O(n \log(n))$

# Parallel Prescan Blelloch



reduction

downsweep

steps:  $2 \log(n)$   
work:  $O(n)$

# Histogram

```
__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}
```

# Histogram (atomics)

```
__global__ void naive_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    d_bins[myBin]++;
}
```

```
__global__ void simple_histo(int *d_bins, const int *d_in, const int BIN_COUNT)
{
    int myId = threadIdx.x + blockDim.x * blockIdx.x;
    int myItem = d_in[myId];
    int myBin = myItem % BIN_COUNT;
    //d_bins[myBin]++;
    atomicAdd(&(d_bins[myBin]), 1);
}
```

# Histogram (reduce-based)

1. Each thread calculates its own histogram

thread0



thread1



thread2




2. Reduce the histogram element by element

# CUDA Streams


## Sequential Version



## Asynchronous Version 1



## Asynchronous Version 2



Time →



# CUDA Streams

```
cudaStream_t mystream;  
cudaError_t result;
```

```
// create and destroy stream  
result = cudaStreamCreate(&mystream);  
result = cudaStreamDestroy(mystream);
```

```
// asynchronously copy from host to device  
result = cudaMemcpyAsync(  
    arr_dev, arr_host, size, cudaMemcpyHostToDevice, mystream  
);
```

```
// launch kernel  
increment<<<gridSize, blockSize, shSize, mystream>>>(arr_dev);
```

# PyCuda

- wrapper of CUDA in Python
- you still write your kernel in C and send it as a string
- try it yourself 😊