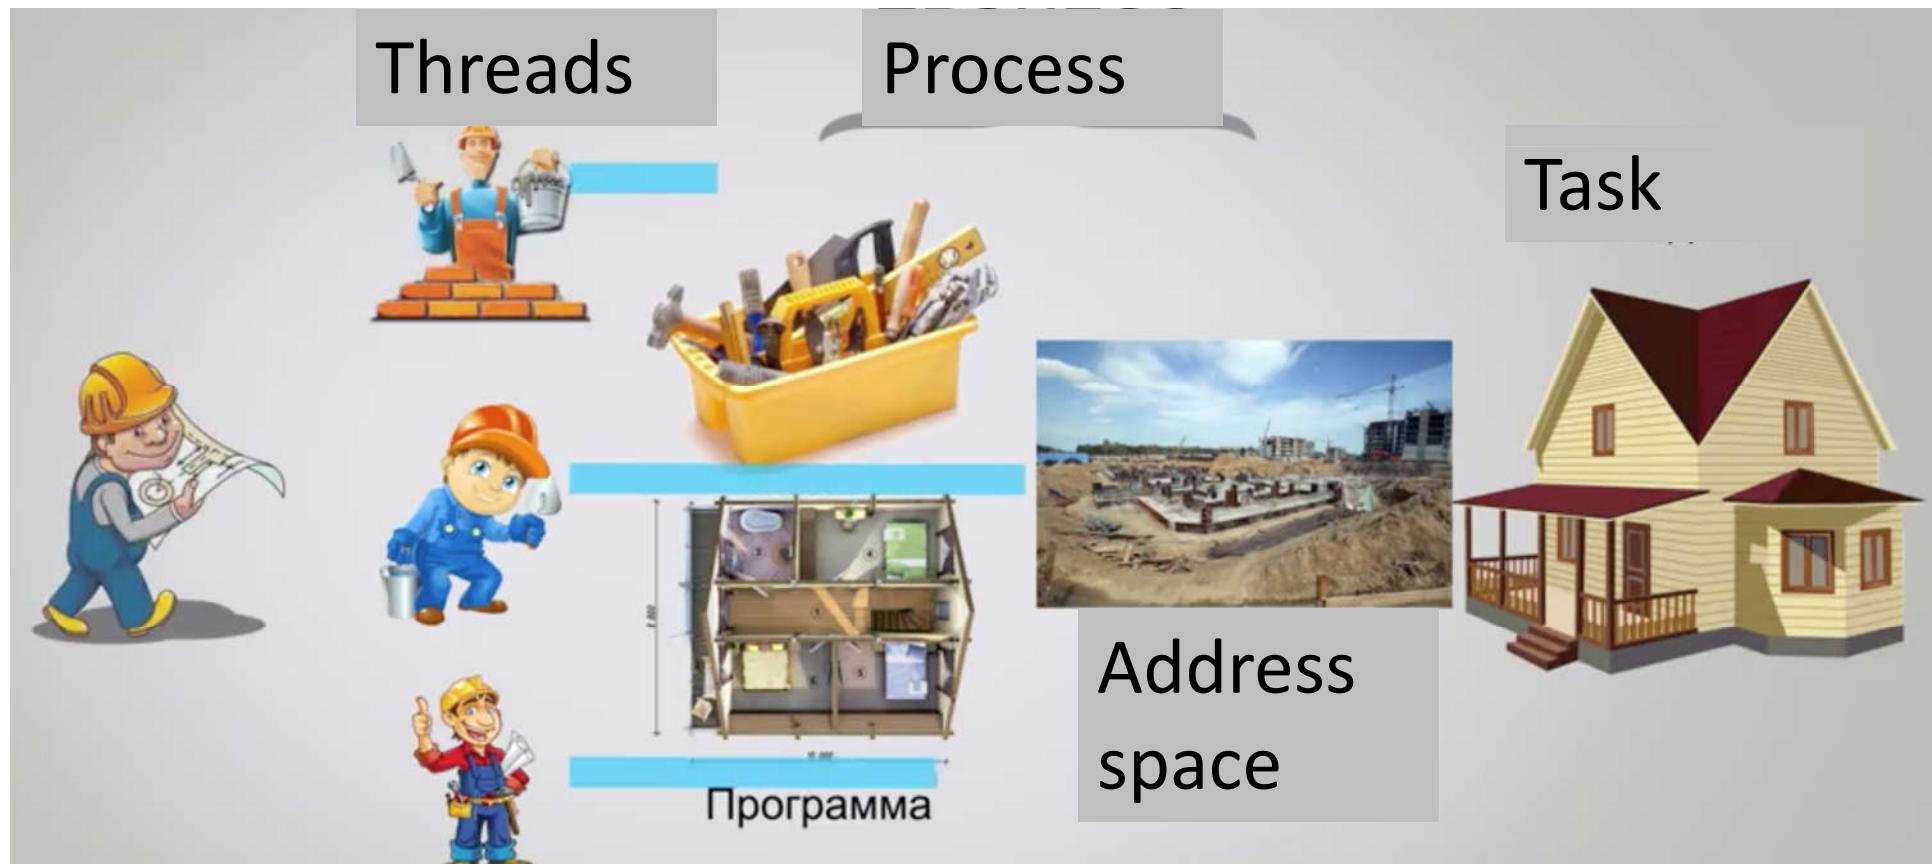
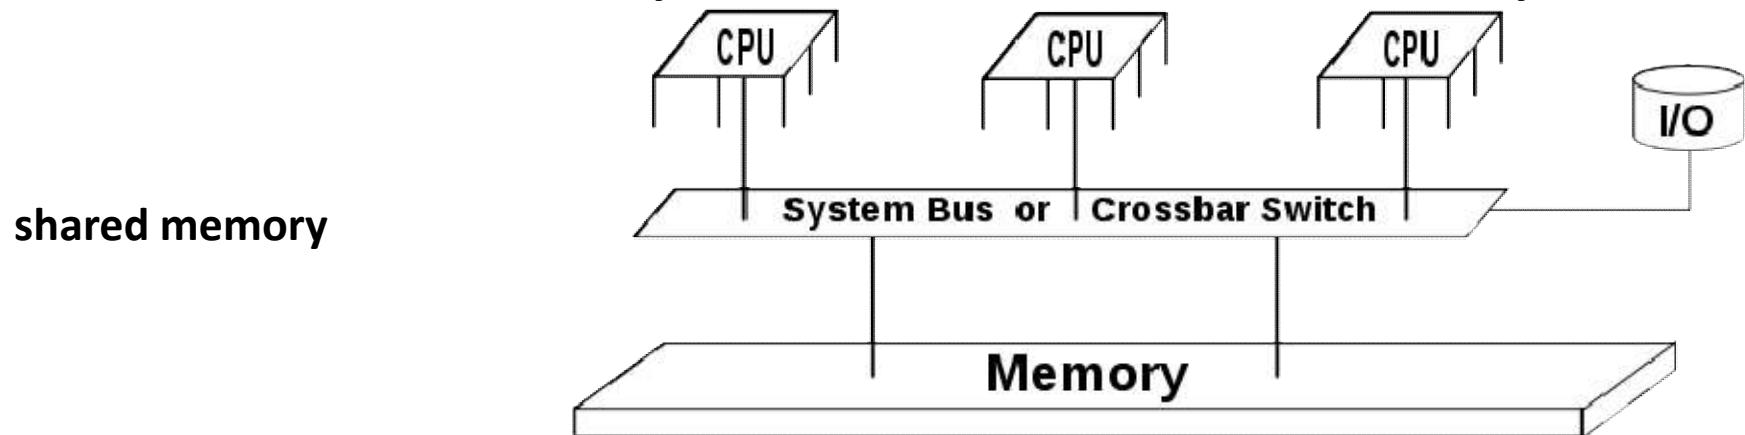


# High Performance Computing

## Term 4 2018/2019

Lecture 3

# Shared memory and distributed memory



# OpenMP

OpenMP is an API for multithreaded, shared memory parallelism.

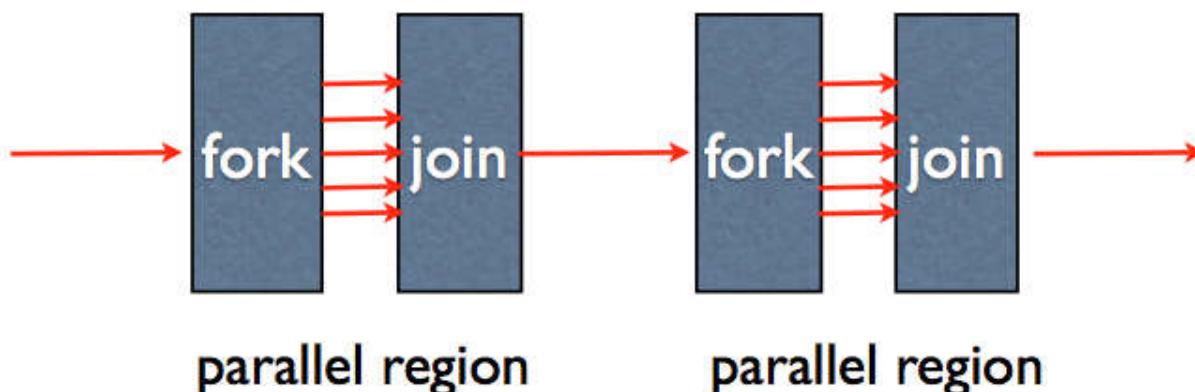
- A set of compiler directives inserted in the source program
  - pragmas in C/C++ (pragma = compiler directive external to prog. lang. for giving additional info., usually non-portable, treated like comments if not understood)
  - (specially written) comments in fortran
- Library functions
- Environment variables

Goal is standardization, ease of use, portability. Allows incremental approach. Significant parallelism possible with just 3 or 4 directives.

# Fork-Join concept

Explicit programmer control of parallelization using fork-join model of parallel execution

- all OpenMP programs begin as single process, the master thread, which executes until a parallel region construct encountered
- FORK: master thread creates team of parallel threads
- JOIN: When threads complete statements in parallel region construct they synchronize and terminate, leaving only the master thread. (similar to fork-join of Pthreads)



# Fork-Join concept

- User inserts directives telling compiler how to execute statements
  - which parts are parallel
  - how to assign code in parallel regions to threads
  - what data is private (local) to threads
  - `#pragma omp` in C and `!$omp` in Fortran
- Compiler generates explicit threaded code
- Rule of thumb: One thread per core (2 or 4 with hyperthreading)
- Dependencies in parallel parts require synchronization between threads

# Simple example

```
#include <omp.h>
#include <stdio.h>

int main() {

#pragma omp parallel
printf("Hello world from thread %d\n",omp_get_thread_num());

return 0;
}
```

**Compile line:**

```
gcc -fopenmp helloworld.c
```

# Setting number of threads

Environment Variables:

setenv OMP\_NUM\_THREADS 2 (cshell)

export OMP\_NUM\_THREADS=2 (bash shell)

Library call:

omp\_set\_num\_threads(2)

```
#include <omp.h>
#include <stdio.h>

int main() {
    omp_set_num_threads(2);

#pragma omp parallel
    printf("Hello world from thread %d\n",omp_get_thread_num());

    return 0;
}
```

# Parallel construct

```
#include <omp.h>

int main() {
    int var1, var2, var3;

    ...serial Code

    #pragma omp parallel private(var1, var2) shared (var3)
    {
        ...parallel section
    }

    ...resume serial code
}
```

# Parallel construct

- When a thread reaches a PARALLEL directive, it becomes the master and has thread number 0.
- All threads execute the same code in the parallel region (Possibly redundant, or use work-sharing constructs to distribute the work)
- There is an implied barrier\* at the end of a parallel section. Only the master thread continues past this point.
- If a thread terminates within a parallel region, all threads will terminate, and the result is undefined.

barrier - all threads wait for each other; no thread proceeds until all threads have reached that point

# Parallel construct

- If program compiled serially, openMP pragmas and comments ignored, stub library for omp library routines
- easy path to parallelization
- One source for both sequential and parallel helps maintenance.

# Work sharing constructs

- work-sharing construct divides work among member threads. Must be dynamically within a parallel region.
- No new threads launched. Construct must be encountered by all threads in the team.
- No implied barrier on entry to a work-sharing construct; Yes at end of construct.

3 types of work-sharing construct (4 in Fortran - array constructs):

- **for loop:** share iterates of `for` loop (“data parallelism”) iterates must be independent
- **sections:** work broken into discrete section, each executed by a thread (“functional parallelism”)
- **single:** section of code executed by one thread only

# For schedule example

```
#include <stdio.h>
#include <omp.h>

#define N 20

int main()
{
    int sum=0;
    int a[N], i;

#pragma omp parallel for
for(i=0;i<N;i++)
{
    a[i]=i;
    printf("iterate i=%3d by thread %3d\n",i, omp_get_thread_num());
}
return 0;
}
```

# OpenMP directives

All directives:

```
#pragma omp directive [clause ...]
                      if (scalar_expression)
                      private (list)
                      shared (list)
                      default (shared | none)
                      firstprivate (list)
                      reduction (operator: list)
                      copyin (list)
                      num_threads (integer-expression)
```

Directives are:

- Case sensitive (not for Fortran)
- Only one directive-name per statement
- Directives apply to at most one succeeding statement, which must be a structured block.
- Continue on succeeding lines with backslash ( " \ " )

# OpenMP for

```
#pragma omp for [clause ...]
    schedule (type [,chunk])
    private (list)
    firstprivate(list)
    lastprivate(list)
    shared (list)
    reduction (operator: list)
    nowait
```

**SCHEDULE:** describes how to divide the loop iterates

- **static** = divided into pieces of size *chunk*, and statically assigned to threads. Default is approx. equal sized chunks (at most 1 per thread)
- **dynamic** = divided into pieces of size *chunk* and dynamically scheduled as requested. Default chunk size 1.
- **guided** = size of chunk decreases over time. (Init. size proportional to the number of unassigned iterations divided by number of threads, decreasing to *chunk size*)
- **runtime**=schedule decision deferred to runtime, set by environment variable OMP\_SCHEDULE.

# OpenMP for example

```
#pragma omp parallel shared(n,a,b,x,y), private(i)
{ // start parallel region

    #pragma omp for nowait
    for (i=0;i<n;i++)
        b[i] = += a[i];

    #pragma omp for nowait
    for (i=0;i<n;i++)
        x[i] = 1./y[i];

} // end parallel region (implied barrier)
```

Spawning tasks is expensive: reuse if possible.  
*nowait* clause: minimize synchronization.

# OpenMP sections example

```
#pragma omp sections [clause ...]
    private (list)
    firstprivate(list)
    lastprivate(list)
    reduction (operator: list)
    nowait
{
    #pragma omp section
        structured block
    #pragma omp section
        structured block
}
```

- implied barrier at the end of a SECTIONS directive, unless a NOWAIT clause used
- for different numbers of threads and SECTIONS some threads get none or more than one
- cannot count on which thread executes which section
- no branching in or out of sections

# OpenMP sections example

```
#pragma omp single [clause ...]
    private (list)
    firstprivate(list)
    nowait
structured block
```

- SINGLE directive says only one thread in the team executes the enclosed code
- useful for code that isn't thread-safe (e.g. I/O)
- rest of threads wait at the end of enclosed code block (unless NOWAIT clause specified)
- no branching in or out of SINGLE block

# OpenMP firstprivate example

What is wrong with this code snippet?

```
#pragma omp parallel for
for (i=0;i<n;i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

# OpenMP firstprivate example

What is wrong with this code snippet?

```
#pragma omp parallel for
for (i=0;i<n;i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

By default, `x` is shared variable (`i` is private).

Could have: Thread 0 set `x` for some `i`.

Thread 1 sets `x` for different `i`.

Thread 0 uses `x` but it is now incorrect.

# OpenMP firstprivate example

Instead use:

```
#pragma omp parallel for private(x)
for (i=0;i<n;i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

What about `i`, `dx`, `y`?

By default `dx`, `n`, `y` shared.

`dx`, `n` used but not changed. `y` changed, but independently for each `i`

# OpenMP firstprivate example

What is wrong with this code?

```
dx = 1/n.;  
#pragma omp parallel for private(x,dx)  
for (i=0;i<n;i++) {  
    x = i*dx  
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);  
}
```

Specifying `dx` private creates a new private variable for each thread, but it is not **initialized**.

**firstprivate** clause creates private variables and initializes to the value from the master thread before the loop.

**lastprivate** copies last value computed by a thread (for  $i=n$ ) to the master thread copy to continue execution.

# Clauses

These clauses not strictly necessary but may be convenient (and may have performance penalties too).

- **lastprivate** private data is undefined after parallel construct. this gives it the value of last iteration (as if sequential) or sections construct (in lexical order).
- **firstprivate** pre-initialize private vars with value of variable with same name before parallel construct.
- **default** (none | shared). In fortran can also have private. Then only need to list exceptions. (none is better habit).
- **nowait** suppress implicit barrier at end of work sharing construct. Cannot ignore at end of parallel region. (But no guarantee that if have 2 `for` loops where second depends on data from first that same threads execute same iterates)

# More clauses

- **if (logical expr)** true = execute parallel region with team of threads; false = run serially (loop too small, too much overhead)
- **reduction** for assoc. and commutative operators compiler helps out; reduction variable is shared by default (no need to specify).

```
#pragma omp parallel for default(none) \
    shared(n,a) \
    reduction(+:sum)
for (i=0;i<n;i++)
    sum += a[i]
/* end of parallel reduction */
```

Also other arithmetic and logical ops., min,max intrinsics in Fortan only.

# Race condition example

```
#include <stdio.h>
#include <omp.h>

int main(){

    int x = 2;

#pragma omp parallel shared(x) num_threads(2)
{
    if (1 == omp_get_thread_num()){
        x = 5;
    }
    else {
        printf("1: Thread %d has x = %d\n",omp_get_thread_num(),x);
    }

#pragma omp barrier

    if (0 == omp_get_thread_num()) {
        printf("2: thread %d has x = %d\n",omp_get_thread_num(),x);
    }
    else {
        printf("3: thread %d has x = %x\n",omp_get_thread_num(),x);
    }

}
return 0;
}
```

# Synchronization in OpenMP

- Implicit **barrier** synchronization at end of parallel region (no explicit support for synch. subset of threads). Can invoke explicitly with `#pragma omp barrier`. All threads must see same sequence of work-sharing and barrier regions .
- **critical sections**: only one thread at a time in critical region with the same name. `#pragma omp critical [ (name) ]`
- **atomic** operation: protects updates to individual memory loc. Only simple expressions allowed. `#pragma omp atomic`
- **locks**: low-level run-time library routines (like mutex vars., semaphores)
- **flush** operation - forces the executing thread to make its values of shared data consistent with shared memory
- **master** (like single but not implied barrier at end), *ordered*, ...

At all these (implicit or explicit ) synchronization points OpenMP ensures that threads have consistent values of shared data.

# Atomic example

```
int sum = 0;  
#pragma omp parallel for shared(n,a,sum)  
{  
    for (i=0; i<n; i++) {  
        #pragma omp atomic  
        sum = sum + a[i];  
    }  
}
```

Better to use a *reduction* clause:

```
int sum = 0;  
#pragma omp parallel for shared(n,a)  \  
    reduction(+:sum)  
{  
    for (i=0; i<n; i++) {  
        sum += a[i];  
    }  
}
```

# Locks

Locks control access to shared resources. Up to implementation to use spin locks (busy waiting) or not.

- Lock variables must be accessed only through locking routines:

`omp_init_lock`    `omp_destroy_lock`

`omp_set_lock`    `omp_unset_lock`    `omp_test_lock`

- In C, lock is a type `omp_lock_t` or `omp_nest_lock_t`  
(In Fortran lock variable is integer)

- initial state of lock is unlocked.

- `omp_set_lock` (`omp_lock_t *lock`) forces calling thread to wait until the specified lock is available.  
(Non-blocking version is `omp_test_lock`)

# Locks

```
#include <stdio.h>
#include <omp.h>

omp_lock_t my_lock;

int main() {
    omp_init_lock(&my_lock);

    #pragma omp parallel
    {
        int tid = omp_get_thread_num( );
        int i, j;

        for (i = 0; i < 5; ++i) {
            omp_set_lock(&my_lock);
            printf("Thread %d - starting locked region\n", tid);

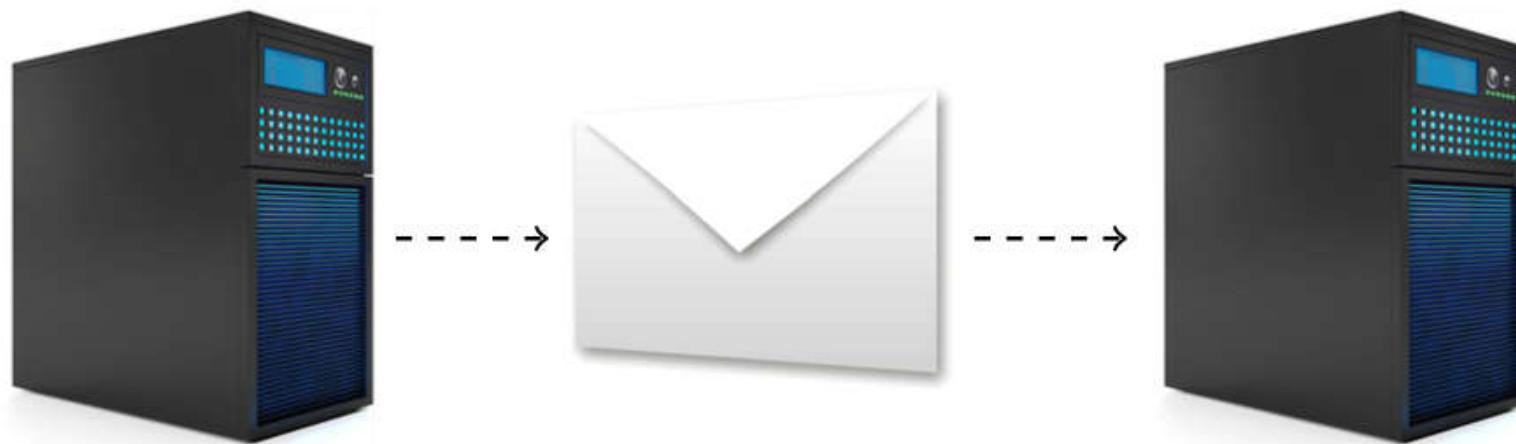
            printf("Thread %d - ending locked region\n", tid);
            omp_unset_lock(&my_lock);
        }
    }

    omp_destroy_lock(&my_lock);
}
```

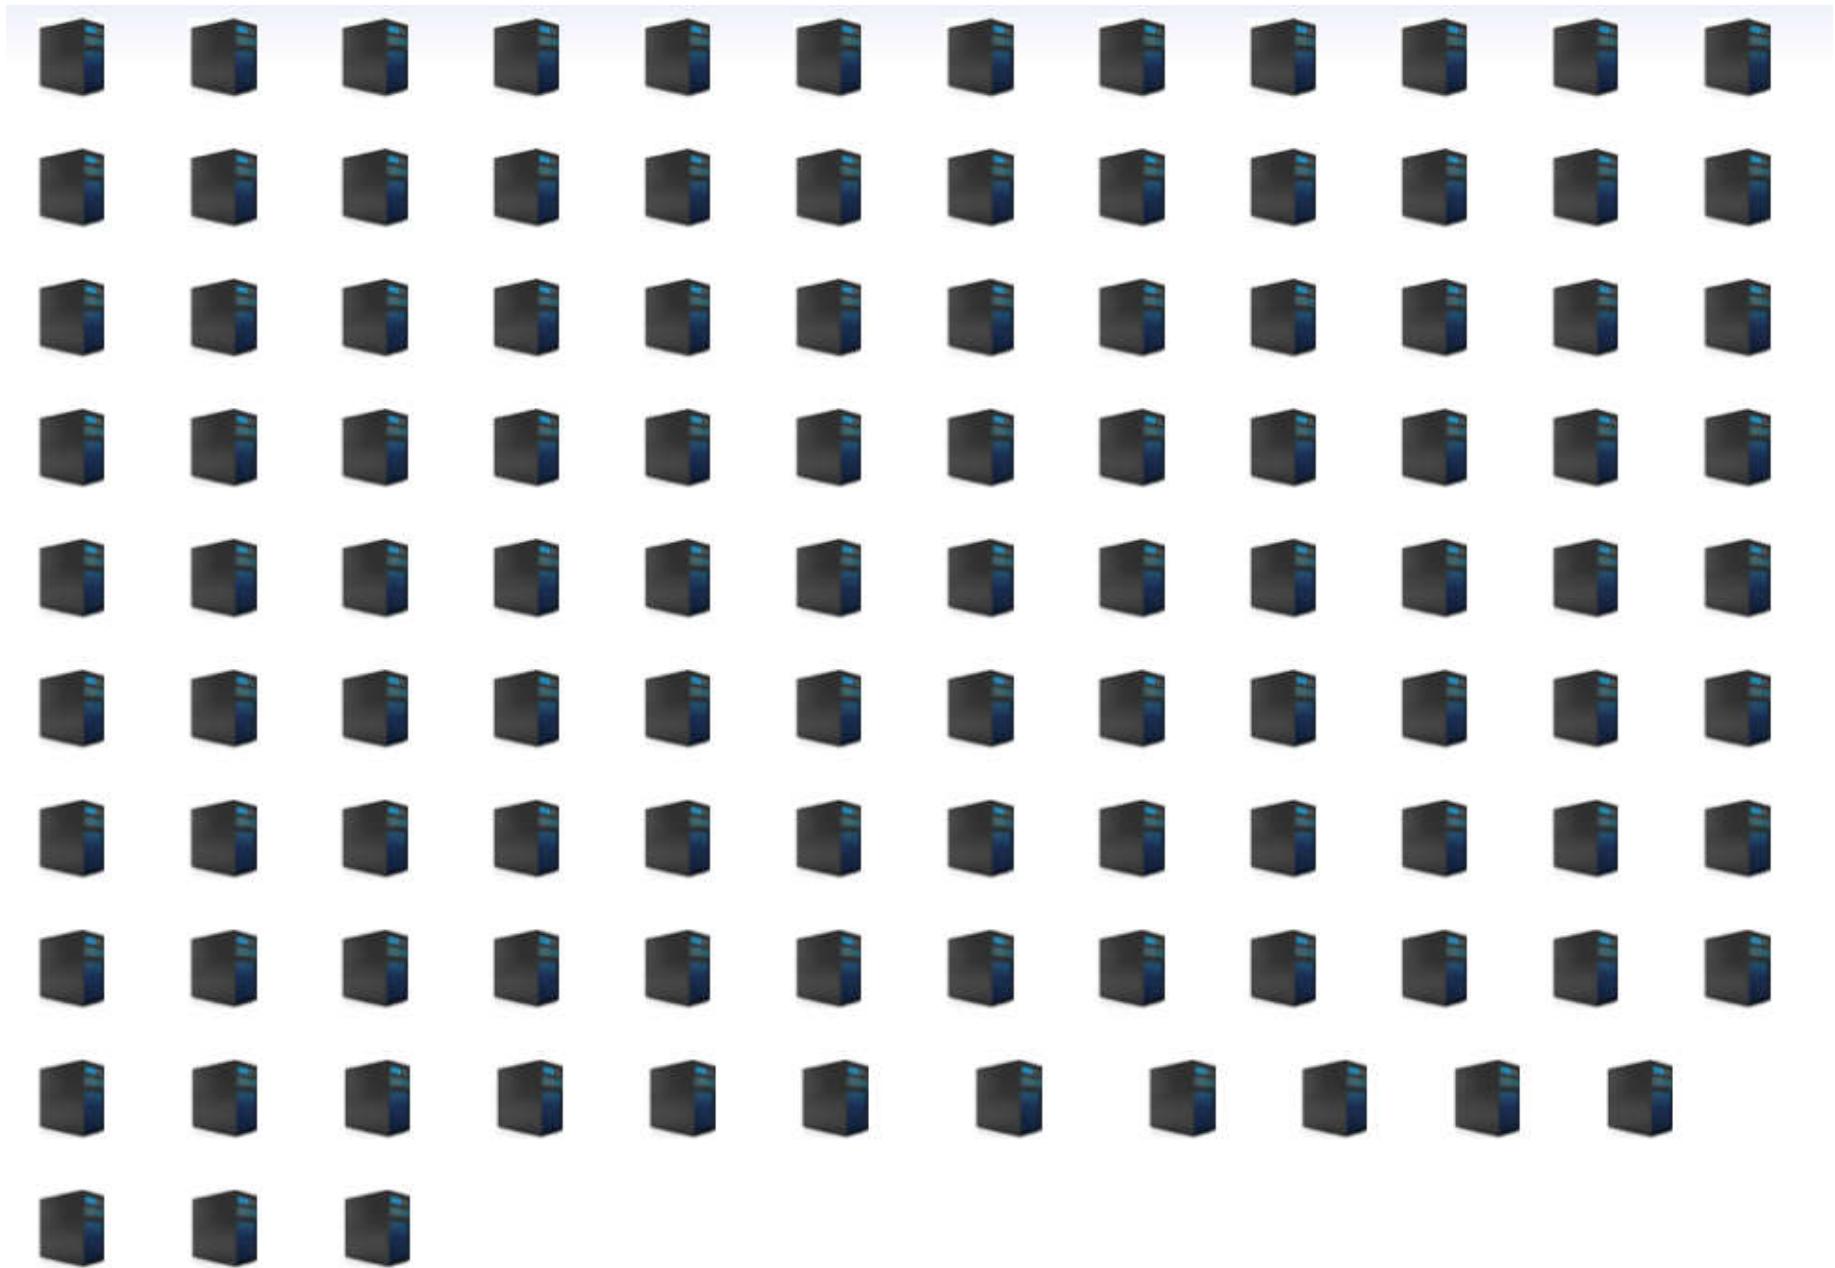


# MPI

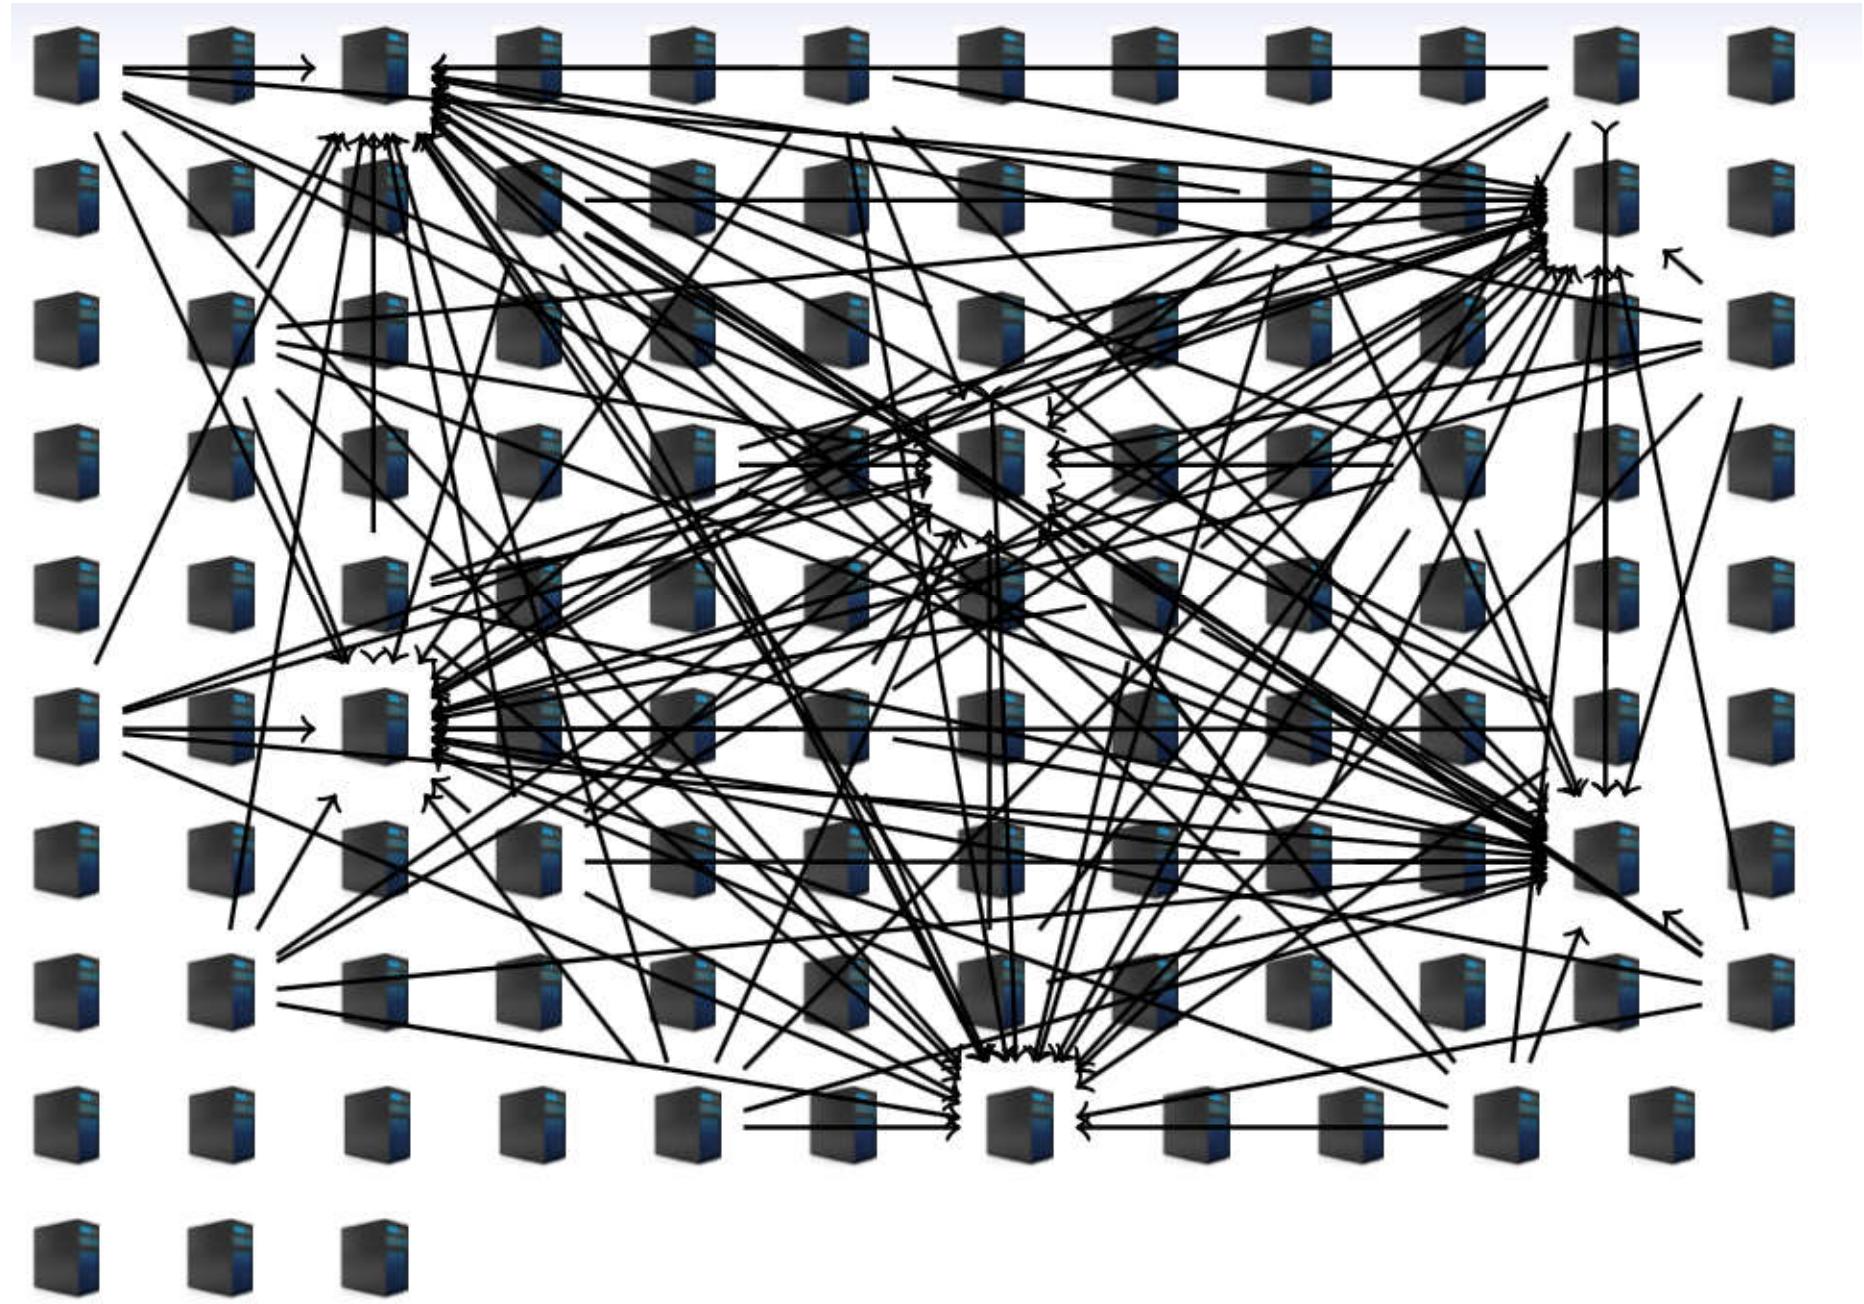
**Message Passing Interface:**



# MPI

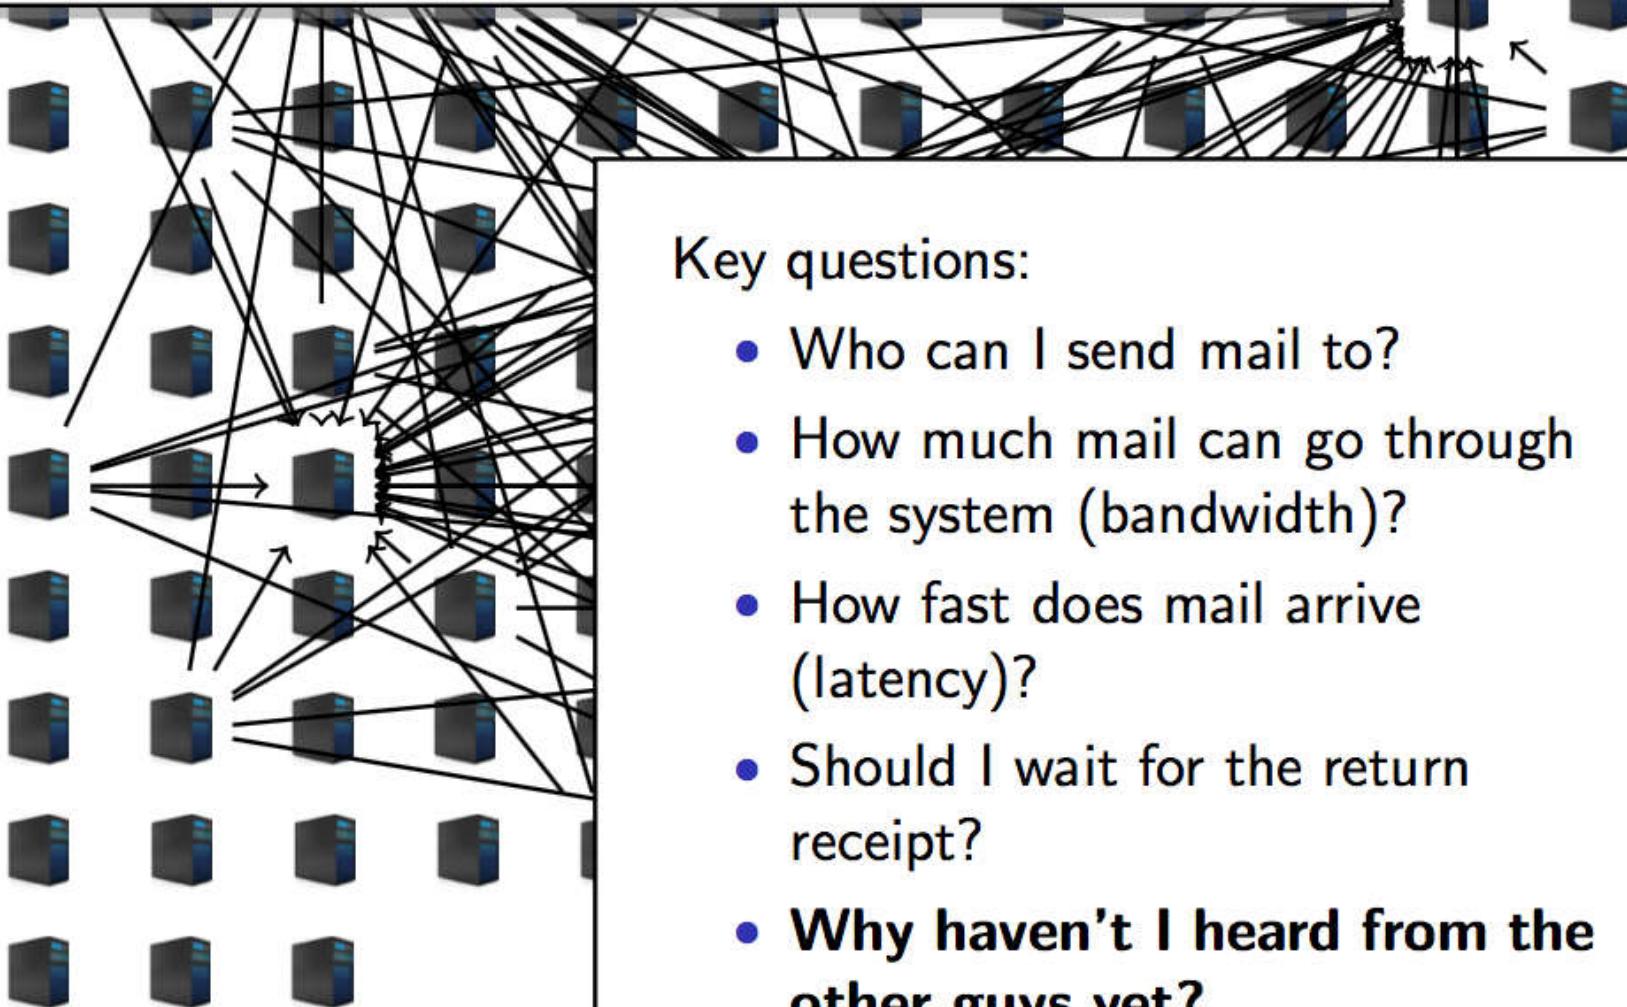


# MPI



# MPI

Not enough throughput? Just buy more computers\*



# MPI hello world

```
# include <mpi.h>
# include <stdio.h>
# include <stdlib.h>

int main(int argc, char ** argv)
{
    int psize;
    int prank;
    MPI_Status status;

    int ierr;

    ierr = MPI_Init(&argc, &argv);
    ierr = MPI_Comm_rank(MPI_COMM_WORLD, &prank);
    ierr = MPI_Comm_size(MPI_COMM_WORLD, &psize);

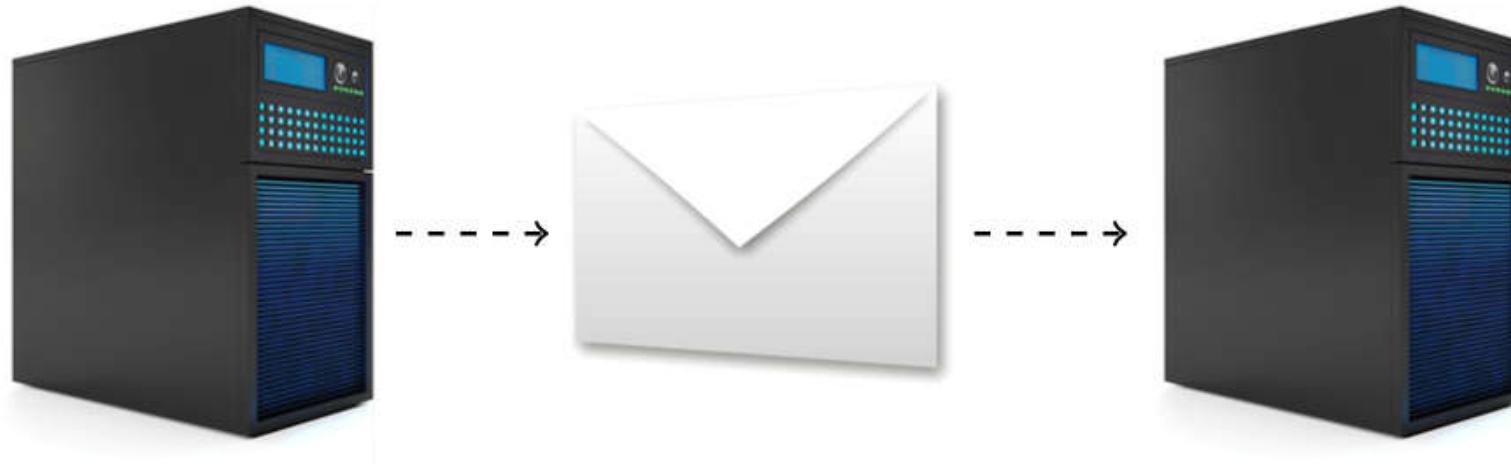
    if (prank == 0)
    {
        printf("The number of processes available is %d\n", psize);
    }

    printf("Hello world from process[%d]\n", prank);

    ierr = MPI_Finalize();
    return 0;
}
```

# MPI. Point-to-point communication

Message Passing Interface:



```
MPI_Send(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int destination,  
    int tag,  
    MPI_Comm communicator)
```

```
MPI_Recv(  
    void* data,  
    int count,  
    MPI_Datatype datatype,  
    int source,  
    int tag,  
    MPI_Comm communicator,  
    MPI_Status* status)
```

# MPI Datatypes

MPI datatype	C equivalent
MPI_SHORT	short int
MPI_INT	int
MPI_LONG	long int
MPI_LONG_LONG	long long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED_LONG_LONG	unsigned long long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	char

# MPI Send&Recv example

```
// Find out rank, size
int world_rank;
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
int world_size;
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

int number;
if (world_rank == 0) {
    number = -1;
    MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
} else if (world_rank == 1) {
    MPI_Recv(&number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
             MPI_STATUS_IGNORE);
    printf("Process 1 received number %d from process 0\n",
           number);
}
```

